When Simulation-Based Sample Supplementation Fails in Cross-Domain Bearing Fault Diagnosis: The Recording-Fingerprint Effect and an Evaluation Protocol That Reveals It

Fangyu Zhang [a,*], Julong He [b]

[a] *College of Mechanical and Aerospace Engineering, Jilin University, Changchun, 130025, China*
[b] *National College for Excellent Engineers, University of Science and Technology Beijing, Beijing, 100083, China*

Abstract

Physics-based simulation, including digital twins, is a popular source of synthetic samples for cross-domain bearing diagnosis when labeled target data are scarce. On the standard CWRU 12 kHz drive-end benchmark it is *structurally* unable to supply class information. Every "fault type × load" cell is a *single recording*, and the class-conditional magnitude spectrum is a stable property of it: a class's spectral envelope recurs across loads at cosine 0.944 but is not predictable from fault physics — a source-domain envelope transfers at 0.924 while one estimated from five *target* labels reaches 0.995. Class identity is thus a per-bearing *recording fingerprint*. A calibrated McFadden–Smith twin produces nine near-identical classes and, trained on twin data alone, reaches macro-F1 0.123 (chance 0.10). The few-shot protocol compounds this: support and test windows share the fingerprint, so cross-recording evaluation lowers macro-F1 by 0.14–0.21 (up to 0.29) and leaves the best augmentation 0.013 above the trivial baseline. Without source data, reproducing the target recording's spectrum — with colored noise or a per-sample phase surrogate — beats the support baseline at every budget of the standard split, while the twin is worse than no augmentation. On

*Corresponding author.
Email addresses: fyzhang9925@mails.jlu.edu.cn (Fangyu Zhang), u202541560@xs.ustb.edu.cn (Julong He)

Paderborn, where several bearings share each class, class-mean spectra remain indistinguishable (cosine 1.000) despite physically defined labels; there same-recording optimism of +0.386 splits into session (+0.161) and specimen (+0.225) terms. The physical component costs 0.26 macro-F1 on CWRU and nothing detectable where the labels follow the fault kinematics.

Keywords: bearing fault diagnosis, benchmark evaluation, cross-domain transfer, digital twin, few-shot learning, simulation-based data augmentation, spectral fingerprint

---

1. Introduction

Deep models for bearing fault diagnosis are notoriously fragile under domain shift, and the scarcity of labeled data on a newly commissioned machine makes transfer learning unavoidable. Two families of remedies dominate the literature: (i) unsupervised domain adaptation, typically adversarial [1] or moment-matching [2], and (ii) the generation of additional labeled samples for the target operating condition, either by physics-based simulation [3, 4, 5] or by generative models [6].

The simulation route is attractive because a physical model is, in principle, transferable [7, 8]: bearing characteristic frequencies scale with shaft speed, so a model of a spalled outer race at 1797 rpm can be re-instantiated at 1730 rpm without any new data. Many recent papers report that augmenting training with simulated faults improves cross-domain accuracy [5, 9, 10, 11, 3, 4].

This paper asks a prior question: *what class-related information can a physical simulator actually supply on the benchmarks we use to test it?* We answer it for the most widely used benchmark, the CWRU 12 kHz drive-end dataset. There the simulator supplies essentially nothing, for a structural reason, and the protocol used to evaluate such methods inflates the results it reports.

1.1. Contributions

- **C1 (diagnosis).** A physically calibrated twin remains 99.99% separable from real data under a domain discriminator, and a classifier trained only on twin data transfers to the real target domain at macro-F1 = 0.123, i.e. chance for ten classes (Section 4).

- **C2 (root cause).** Class identity in CWRU is dominated by a *recording fingerprint*. Per-class spectral envelopes are stable across loads (cosine 0.944 averaged over all load pairs), but the source-domain envelope predicts the target-domain envelope only to 0.924. The twin reproduces neither (Section 5).

- **C3 (protocol).** The few-shot evaluation protocol draws support and test windows from the same recording, so both share the fingerprint. Under a cross-recording protocol the reported macro-F1 drops by 0.14–0.21 (average over four support domains; up to 0.29), falling to 0.02 when the three fault sizes are pooled into one class, and fingerprint augmentation ends within 0.015 of training on the labeled support set alone, while the phase surrogate falls 0.06 behind it (Section 6). We recommend it for any few-shot result on this benchmark.

- **C4 (constructive).** The same mechanism says what an augmentation has to supply: the target recording's spectrum, not the fault physics. The class envelope is easy to learn from a few labels (cosine 0.995 at $k = 5$), and coloring noise with it beats training on the support set at every budget (+0.077, +0.037, +0.036 macro-F1 at $k = 1, 3, 5$; 17–19 of 20 paired runs). A per-sample phase surrogate drops the class average and keeps each support window's own spectrum, and it matches the envelope: fidelity to the recording is what carries the gain. The physics twin is worse than no augmentation at every budget, and its impulse train adds nothing to the envelope (Sections 8 and 9).

- **C5 (side finding).** Load 3 of CWRU is an outlier operating condition: a label-free model transfers among loads 0–2 at 0.87–1.00 but reaches 0.63–0.68 on load 3, and labeled support from the other loads transfers into load 3 at 0.69–0.80. Conclusions drawn from the popular S0→S3 direction alone do not generalize (Section 6).

- **C6 (cross-benchmark).** On Paderborn, where several independent bearings share each damage class, class-mean magnitude spectra of healthy, inner-ring and outer-ring bearings are identical (cosine 1.000; CWRU: 0.626), and remain so after the dominant rig tone is filtered out (margin +0.003). The CWRU class separation is therefore a property of the individual specimens, not of the classes. Same-recording evaluation

is optimistic by +0.386 macro-F1 there; because each Paderborn bearing was measured 20 times, that optimism can be split by re-evaluating on a different *recording of the same bearings*, which gives +0.161 for the session and a further +0.225 for the change of specimen, and it survives removal of the rig tone that holds 93% of the spectral energy (+0.321 remains). A twin whose physics exactly matches the class definition is still the worst augmentation tried, and the physical component that costs 0.26 macro-F1 on CWRU does no detectable harm here — an asymmetry the mechanism explains (Section 7).

2. Related Work

Cross-domain and few-shot bearing diagnosis. The problem setting we study is standard: a model trained under one operating condition must be reused under another, and labels are scarce on the target [12, 13]. The dominant response is unsupervised domain adaptation, usually adversarial, which aligns source and target feature distributions without target labels [1, 14, 15]. A complementary line attacks label scarcity directly through meta-learning and metric-based few-shot schemes [16], while the backbone architectures themselves are now well surveyed [17]. Both are typically evaluated on CWRU or a handful of sibling benchmarks, and both depend on the support set being representative of the target test distribution. We show that under the usual window-level split this assumption is violated in a specific and quantifiable way.

Simulation and digital-twin sample supplementation. Because a physical model transfers by construction, generating labeled fault samples by simulation is attractive precisely when target labels are missing. Recent work builds bearing dynamics models and uses the resulting *twin data* as a labeled source: partial domain adaptation driven by a digital twin [10], discriminative graph learning that transfers knowledge from simulated to real signals [9], multimode collaborative transfer learning on twin samples [5], and digital-twin-assisted domain adversarial frameworks for machinery such as centrifugal pumps [11]. A recurring theme in this literature is the simulation-to-reality gap, and much of the machinery — physics-enhanced translation between simulated and measured signals [3], generative alignment of the two distributions [4] — exists to close it. The same gap motivates domain randomization in robotics [18], where the simulator is deliberately

not made faithful and the burden shifts entirely onto the augmentation distribution.

We ask a question upstream of that gap: what class-discriminative information the simulated data can carry at all, given the structure of the benchmark. On CWRU the answer is none, and closing the gap cannot change it. The class-conditional spectrum there is a property of the individual recording rather than of the fault, so a simulator with perfect fidelity to the physics still cannot supply the contrast that separates the classes.

Surrogate data. Our phase-randomization control follows the standard surrogate-data procedure for testing whether a statistic depends on structure beyond the power spectrum [19, 20]. We use it in the reverse direction from the usual application: rather than asking whether a signal is nonlinear, we use phase-randomized surrogates as a deliberately *structure-free* augmentation baseline, to separate the contribution of the spectral envelope from that of any temporal structure.

3. Data, Taxonomy and Protocol

3.1. CWRU and the ten-class taxonomy

We use the CWRU 12 kHz drive-end recordings [21]. The conventional four-class taxonomy (normal, ball, inner race, outer race) is saturated on cross-load transfer — a source-only model reaches 100% accuracy in all twelve transfer directions — so we adopt a ten-class taxonomy: $\{\textsc{Normal}\} \cup \{B, IR, OR\} \times \{0.007'', 0.014'', 0.021''\}$, using the official files for the four loads $\{0, 1, 2, 3\}$ (1730–1797 rpm). The ten classes are the dataset's own fault-size labels (three sizes for each of ball, inner and outer race); the finer taxonomy is the one in which the methods still differ, since the four-class one saturates. Windows of 1024 samples with stride 512 are z-scored individually.[1]

3.2. The few-shot split

For a target load, each class is split in time order: the last 40% of windows form the held-out test set; from the first 60%, k windows per class are drawn

[1] Every spectral scalar reported in this paper is reproduced from the raw data by one script released with the code (`tools/probe_spectral_numbers.py`), and every table entry by `tools/verify_paper_numbers.py`.

(with the seed) as the labeled *support* set, and the rest form an unlabeled *pool*. The support and test windows come from the *same recording*. Section 6 shows how much of the reported score that choice accounts for.

3.3. Classifier and training

Every configuration is trained with the same model, so that each comparison isolates the training set rather than the architecture. The input is the raw time-domain window, z-scored sample-wise, and no spectral transform is applied before the classifier: the spectral statistics of Sections 4 and 5 are diagnostics on the data, not features of the model. The backbone is a one-dimensional residual network with a 7-tap convolution (stride 2, 32 channels), batch normalization and ReLU, three residual blocks of two 3-tap convolutions each (32, 64 and 128 channels, stride 2 from the second block), global average pooling, and a linear projection to 64 features. One linear layer maps those features to class scores, and the domain discriminator used by the adversarial methods is a two-hidden-layer MLP of width 64 on the same features. The backbone has 118 368 parameters and the model 127 403 including the discriminator. Training minimizes cross-entropy, plus a binary domain loss where a discriminator is used, with Adam at a learning rate of 10^{-3}, batch size 64 and 60 epochs, 200 for the Paderborn comparisons of Section 7; synthetic sets are generated with 400 windows per class. One seed controls the few-shot draw, the synthetic sample generation and the weight initialization jointly.

Bitwise reproducibility is not available on this hardware. Enabling the deterministic cuDNN path makes a seed fail to converge — in our check seed 1's training-set macro-F1 fell from 1.000 to 0.52 — so it stays off, and reproducibility rests on statistics over repeated runs instead. Where a comparison is close we repeat the whole configuration: Tables 1, 3 and 9 report 20 runs (ten seeds, two repetitions each), and the Paderborn tables report ten seeds.

4. Diagnosis I: What the Twin Supplies

4.1. The generator

We implement a McFadden–Smith style impulse-response twin [22] with SKF 6205-2RS geometry, whose ball-pass frequencies are computed from the shaft speed. The generator is calibrated against real CWRU spectra: the

resonance range, the resonance bandwidth per defect size, the background-noise coloring, and the impulse-to-harmonic power ratio were all fitted so that synthetic band-energy shares match measured ones. In particular the synthetic peak envelope spectrum lands on the correct BPFO, so the twin reproduces the fault kinematics that the benchmark labels describe. Resonance frequency, impulse-to-background ratio, and harmonic amplitudes and phases are redrawn for every synthetic window, and the impulse times carry random jitter, so the generator used here is a distribution of twins rather than one tuned point in its parameter space.

The generator enters the model the way the cited works use theirs: as the labeled source of a domain-adversarial configuration adapted to the target load (DT_DANN, Table 8). The failure diagnosed below therefore belongs to that template: a simulated source aligned adversarially with the target.

4.2. Twin data are a different domain, and carry no class information

Training only on twin data and testing on the real target load (Table 1) gives macro-F1 = 0.123±0.030 against a chance level of 0.10. A linear domain discriminator on backbone features separates twin from real with accuracy 0.9999. The same table shows the ordering that the rest of the paper is about: the target-calibrated envelope reaches 0.910, and every attempt to put physical structure back on top of it *lowers* the score.

What moves the score. Table 1 separates two ways of improving a synthetic source: greater fidelity to the physics, and more information about the target. Relative to the twin alone (0.123), a measured source-domain envelope reaches 0.587 with white-noise excitation and 0.491 with the impulse train, so adding physics lowers it. Replacing the source envelope with the *target* envelope gives 0.651 with the impulse train and 0.910 with white noise. These two generators are oracles — their input is the class-conditional envelope of the target recording, which no simulator can compute without target data — and even they do better without physics. What the score responds to is how much of the target recording the generator is given, not how faithful its physics is: a more accurate impulse model moves the twin toward the measured source envelope (0.587), not toward the target-envelope value (0.910).

Table 1: Training only on a synthetic source, testing on the real target load (load 3, ten classes). Each entry is 20 runs — ten seeds, the whole configuration repeated twice — and the ± is the standard deviation over those 20 runs. "Disc." is a domain discriminator trained on backbone features to separate the synthetic source from real target data (0.5 = indistinguishable). Every synthetic set is drawn with a randomized resonance frequency (2850–3550 Hz), impulse power and harmonic amplitudes and phases, and with jittered impulse timing, so each row averages over generator parameters as well as seeds. The two repetitions of the target-fingerprint row differ by 0.033 in the mean of ten seeds, and a single ten-seed draw would understate that row's spread by a factor of two.

Synthetic source	macro-F1	Disc.
Physics twin (source speed)	0.123 ± 0.030	0.9999
Source fingerprint + white noise	0.587 ± 0.045	0.9972
Source fingerprint + impulse train	0.491 ± 0.041	0.9994
Target fingerprint + impulse train	0.651 ± 0.056	0.9786
Target fingerprint + impulse (comb-free)	0.524 ± 0.061	0.9999
Target fingerprint + white noise	$\mathbf{0.910 \pm 0.077}$	**0.9337**

5. Diagnosis II: the Recording Fingerprint

5.1. Class differences are stable across loads

Taking the smoothed magnitude envelope $A_c(f)$ as each class's spectral signature (Section 8), the signature of load 0 correlates with those of loads 1–3 at cosine 0.954, 0.939 and 0.924 respectively (mean 0.939); averaged over all six load pairs the value is 0.944. Class differences therefore survive a change of load and are a stable property of the *recording*, which is why real-data methods work well here. Table 4 reports a separate, unsmoothed statistic (Section 7), which yields systematically lower cosines.

5.2. But they are not predictable from fault physics

The nine synthetic fault classes have near-identical mean spectra (Fig. 2(b)): all occupy 48–64% (mean 55%) of their energy in the 2800–3300 Hz band, whereas the nine real classes at load 0 spread from 9% (OR014) to 67% (IR021), mean 32%. Table 2 quantifies this with the mean pairwise cosine distance between class-mean spectra: real data 0.374, twin 0.167.

Table 2 also exposes a subtlety that we return to in Section 9: the physics impulse train *increases* inter-class distance relative to the fingerprint-only

Table 2: Mean pairwise cosine distance between class-mean magnitude spectra (higher = classes better separated). Target load 3.

Source	Mean inter-class distance
Real target (load 3)	0.374
Real source (load 0)	0.374
Physics twin	0.167
Target fingerprint + white noise	0.315
Target fingerprint + impulse train	0.346
Target fingerprint (comb-free) + imp.	0.319

condition, yet it transfers worse. The physical structure does not destroy class information; it pushes the sample distribution away from the real one.

5.3. Why physics cannot close the gap

Each CWRU fault file was recorded with a *different* damaged bearing, and the same bearing was re-used across loads for a given file. "What this class looks like" is therefore a property of that particular bearing and recording session. A simulator cannot know which bearing a target recording used. The fingerprint is estimable from target labels (Section 8) but not derivable from fault mechanics. In the terminology of the shortcut-learning literature [23], the classifier learns the recording as the shortcut to the label; here that shortcut is the only signal present.

6. Evaluation Protocol: Same-Recording Leakage

The support and test windows of the standard few-shot split come from one recording, so they share the fingerprint. This is a leakage of the kind that has been catalogued as a general threat to machine-learning-based science [24]: the split is drawn at the window level while the statistical unit that carries the label is the recording, so the test set is not independent of the training set in the sense that matters. To quantify the resulting optimism we evaluate every configuration twice: on the held-out windows of the *same* load, and on the *other* three loads, which are different recordings. Table 3 reports the result for $k = 5$.

Five conclusions follow.

Table 3: Same-recording versus cross-recording evaluation ($k = 5$, macro-F1). Each cell is 20 runs — ten seeds, the whole configuration repeated twice. Each block trains on the labeled support of one load and evaluates on its own held-out windows (*same*) and on the other three loads (*cross*). "Surrogate" = Fourier phase randomization of each support sample; "true aug" = 400 real labeled samples per class drawn from the unlabeled pool.

Support load	Training configuration	Same	Cross mean	Cross (to load 3)	Optimism
load 0	support only	0.989	0.880	0.718	+0.109
	+ fingerprint aug.	0.987	0.851	0.730	+0.136
	+ phase surrogate	0.979	0.747	0.640	+0.232
	+ true samples	0.976	0.837	0.681	+0.139
load 1	support only	0.988	0.869	0.691	+0.119
	+ fingerprint aug.	0.993	0.849	0.694	+0.144
	+ phase surrogate	0.987	0.806	0.654	+0.181
	+ true samples	1.000	0.907	0.748	+0.093
load 2	support only	0.998	0.882	0.798	+0.116
	+ fingerprint aug.	0.990	0.884	0.805	+0.106
	+ phase surrogate	0.993	0.848	0.770	+0.144
	+ true samples	0.969	0.878	0.844	+0.091
load 3	support only	0.993	0.732	—	+0.261
	fingerprint aug.	0.995	**0.831**	—	+0.164
	+ phase surrogate	0.994	0.705	—	+0.289
	+ true samples	1.000	0.719	—	+0.281
Mean over the four support loads					
	support only	0.992	0.840		+0.152
	fingerprint aug.	0.991	0.854		+0.138
	phase surrogate	0.988	0.777		+0.211
	true samples	0.986	0.835		+0.151

Same-recording evaluation is optimistic by 0.14–0.21. Support-only training scores 0.988–0.998 on its own held-out windows but only 0.732–0.882 on other recordings (Fig. 1). Every training configuration in the table shows the same gap, and how large it is depends on how the classes are defined (last paragraph of this subsection). A few-shot result on this benchmark that uses the standard split contains this optimism unless a cross-recording control is reported alongside it.

Augmentation does not buy cross-recording transfer. Averaged over the four support loads, cross-recording macro-F1 is 0.840 for support only, 0.854 for fingerprint augmentation, 0.835 for adding 400 real labeled samples per class, and 0.777 for phase surrogates. Only the surrogate separates from the base-

line by more than noise, and it separates downward. Fingerprint augmentation's +0.013 average is the mean of −0.029, −0.020, +0.002 and +0.099 over the four support loads: two losses, a tie and one gain. Adding large amounts of *real* data from the same operating condition does not improve cross-recording generalization either: the bottleneck is the operating condition, not the sample count.

Fingerprint augmentation is a regularizer. Its advantage appears only when the support recording is atypical. With load 3 as the support load it reaches 0.831 against 0.732 for support only, a gain of 0.099 over twenty runs; with loads 0 and 1 as the support load it is worse by 0.02–0.03 (0.851 against 0.880 and 0.849 against 0.869); with load 2 it is flat (0.884 against 0.882). The mechanism is that averaging the k support spectra into a class-level model suppresses recording-specific idiosyncrasies, which costs accuracy when the support recording is representative and recovers it when the support recording is not.

Load 3 is an outlier condition. A purely source-trained model, using no target labels at all, transfers among loads 0–2 at 0.87–1.00 but reaches 0.63–0.68 when load 3 is the test recording; conversely, labeled support from loads 0–2 transfers into load 3 at 0.69–0.80. The widely used S0→S3 direction is hard mainly because load 3 is peculiar; results reported on that direction alone should not be generalized.

The inflation requires the class to be a single recording. The ten classes are three fault types at three fault sizes, and within a load each class is a single recording. Pool the three sizes into one class — so each fault class spans three recordings — and the same-recording advantage largely disappears: support-only training scores 0.874 on its own held-out windows and 0.854 on the other loads, an optimism of +0.020 against +0.152 for the ten-class taxonomy (ten seeds). Per support load the gap is −0.020, −0.004, +0.015 and +0.088; only load 3 retains one. The cross-recording score is essentially unchanged (0.854 against 0.840); what falls is the same-recording score, from 0.992 to 0.874. Four classes at a single fault size are at ceiling on both evaluations (cross-recording 0.93–1.00), so that control is uninformative. The pooled control is the one that speaks to the mechanism: once a class spans several recordings, a random window split is no longer a same-recording split, and the advantage goes with it.

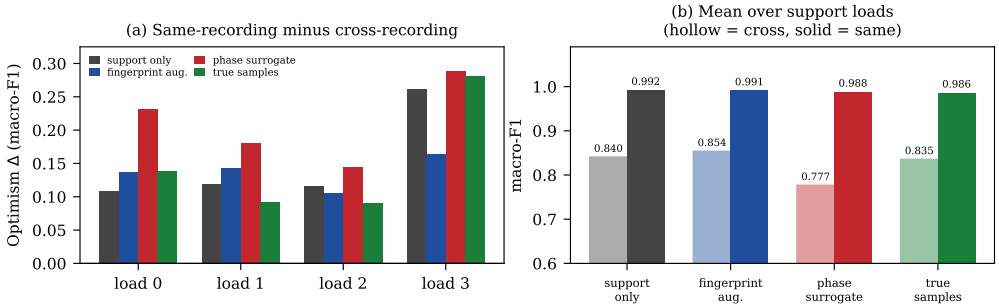


Figure 1: Same-recording versus cross-recording evaluation ($k = 5$, twenty runs: ten seeds, two repetitions each). (a) The optimism Δ of each training configuration, per support load: every bar is positive, i.e. every configuration looks better on its own held-out windows than on other recordings, and the effect is largest exactly where load 3 is involved — the condition used by the most widely reported transfer direction. (b) Averaged over the four support loads. Hollow bars are cross-recording, solid bars same-recording. Once the evaluation is cross-recording the augmentations cluster on the trivial baseline: adding 400 *real* labeled samples per class scores 0.835 against 0.840 for support only, fingerprint augmentation 0.854, and the phase surrogate 0.777.

7. Cross-Benchmark Control: Paderborn

CWRU cannot falsify the fingerprint explanation, because it does not contain the control that would do so: every class has exactly *one* specimen. We therefore repeat the essential measurements on the Paderborn (KAt) dataset [25], in which several independent bearings share each damage class, measured under four operating conditions at 64 kHz.

We use 21 bearings: 6 healthy (K001–K006), 9 artificially damaged (KA series) and 6 damaged by accelerated lifetime tests (KI series). Classes are defined by *damage location*, read from the damage profile that ships with each bearing (outer ring OR, inner ring IR, healthy) — i.e. exactly the kinematic distinction that a fault model can express. Windows are 4096 samples (64 ms). At 64 kHz a 1024-sample window spans only 16 ms — barely more than one BPFO period at 1500 rpm (76 Hz, 13 ms), so a window rarely holds more than one impact — and with 1024-sample windows all bearings produce indistinguishable spectra, whereas the 4096-sample window holds about five periods.

In the subset we downloaded, every KA bearing is outer-ring and every KI bearing is inner-ring, so damage location and damage origin (artificial versus run-to-failure) are confounded. The finding below is an *absence* of

Table 4: Class-mean magnitude spectra, identical statistic on both benchmarks: the *unsmoothed* mean magnitude spectrum of each class, so the two benchmarks are directly comparable. It is distinct from the smoothed envelope $A_c(f)$ of Section 8 used elsewhere, which yields higher cosines. "Within" = different specimens, same class; "between" = different specimens, different class. CWRU has one specimen per class, so only the between-class numbers exist there; its "same class across loads" row is the closest analogue of a within-class measurement. The last two rows remove the rig tone of Section 7 by high-pass filtering at 2 kHz, the control for a dominant rig line.

Benchmark	Quantity	Cosine	Margin
CWRU (one specimen per class)			
	between class, load 0	0.626	—
	between class, load 3	0.626	—
	same class, load 0 vs 1	0.870	—
*Paderborn (*21 *specimens)*			
	within class	1.000	
	between class	1.000	+0.000
Paderborn, high-pass 2 *kHz*			
	within class	0.992	
	between class	0.989	+0.003

class-conditional spectral structure and is unaffected, but the data cannot support a separate statement about outer versus inner race.

7.1. Class-conditional spectra are a rig property, not a class property

Table 4 compares the two benchmarks with one identical statistic: the cosine similarity between class-mean magnitude spectra, shown directly in Fig. 2.

On CWRU the ten classes are strongly separated in the magnitude spectrum (cosine 0.626). On Paderborn, where a class contains several independently manufactured bearings, averaging over specimens leaves *no* separation at all: within-class and between-class cosines are both 1.000 to three decimals. The spectra are dominated by a rig tone: at 1500 rpm the largest spectral line is at 100 Hz — exactly four times the shaft rate — in *every one* of the 21 bearings (measured on the full records it falls in 99.6–100.1 Hz for all of them), healthy, outer-ring and inner-ring alike. In a Hann-windowed 4096-point spectrum that line lands in the bin at 93.75 Hz and *that single bin* carries 54% of the total spectral energy (53.97–54.31% across the 21

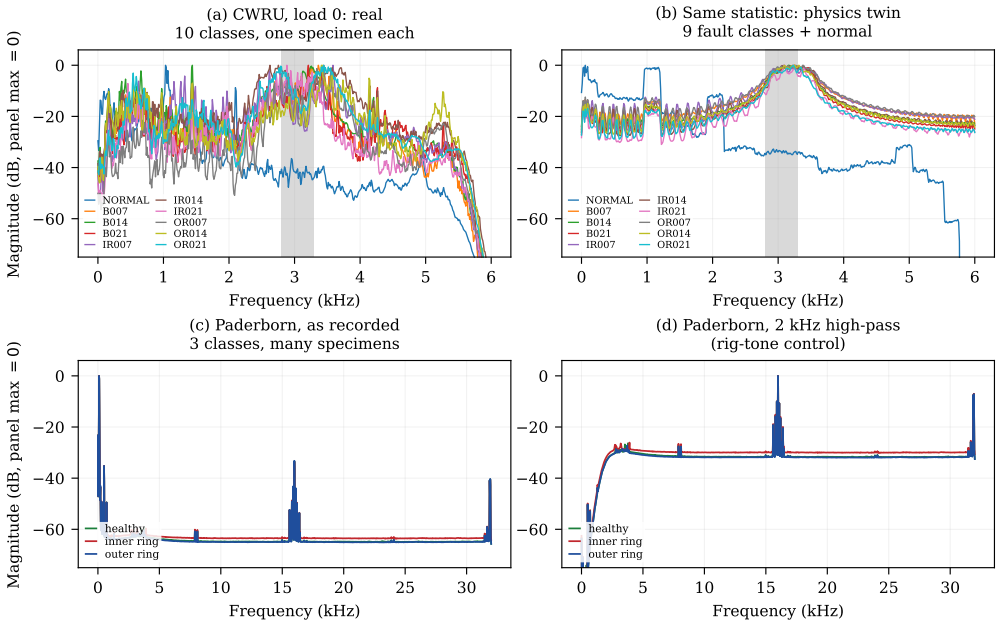


Figure 2: Class-mean magnitude spectra under the identical, unsmoothed statistic of Table 4. Each panel is normalized to its own maximum (0 dB); all four share the same −75 dB floor, the same linear frequency axis and the same per-class colors, and the shaded span marks the 2800–3300 Hz band whose energy fraction is quoted in Section 5. (a) CWRU, ten classes, one physical specimen each: visibly separated. (b) The physics twin, same statistic: the nine fault classes collapse into one band. Their median pairwise separation is 2.5 dB (minimum 0.9 dB), against 6.7 dB (minimum 3.7 dB) for the nine real classes, and across frequency the nine twin curves span a median of 6.5 dB against 18.8 dB. This is the whole negative result in one panel — the simulator produces nine faults whose mean spectra are interchangeable, so no classifier trained on them can carry the contrast that separates the real classes. (c) Paderborn as recorded: the 100 Hz rig line of Section 7 lands in the 93.75 Hz bin (with leakage into its neighbor at 109 Hz) and carries essentially all the energy, and the three damage classes are indistinguishable because everything else lies 60 dB below it. (d) The same recordings after a 2 kHz high-pass removes that line, so the comparison no longer rests on it. The 16 kHz rig resonance now dominates, yet the three class means still track each other to within a few dB across the whole band: the absence of class structure survives removal of the tone. Compare (d) with (a), where the same statistic separates the real classes by 10–30 dB.

Table 5: Same-recording versus bearing-disjoint evaluation on Paderborn (condition N15_M07_F10, three classes, 5 training bearings per class, 200 epochs, ten seeds, macro-F1). The training set is files 1–4 of the training bearings; the same-recording test set is its held-out 30% of windows. All ten runs are included, including the one that did not converge (see the text below).

	Same recording	Bearing-disjoint	Optimism
Mean over 10 seeds	0.891	0.505	+0.386
Std. deviation	0.120	0.161	
CWRU, for comparison (mean over four support loads)			
Support only	0.992	0.840	+0.152

bearings); the line together with two neighboring bins on each side carries 99.5–99.8%.

Because that single line holds almost all the energy, the coincidence of class means could in principle be an artifact of it. After high-pass filtering at 2 kHz the within-class cosine falls to 0.992 and the between-class cosine to 0.989; the margin moves from 0.000 to +0.003. Removing the rig tone does not reveal a hidden class signature — there is none in the magnitude spectrum — the opposite of what happens on CWRU, where the between-class cosine stays at 0.62–0.63 in every load. Filtering before *training* rather than before the statistic gives the same answer; Section 7.4 reports that control.

The two facts together identify the mechanism. The class separation that makes CWRU look easy is not a property of the fault classes; it is a property of the individual specimens and recordings that happen to carry the class labels. Where several specimens share a label, the separation disappears.

7.2. Same-recording leakage reproduces on a second benchmark

A second benchmark tests whether the optimism of Section 6 is specific to CWRU. We train one model per seed on five bearings per class and evaluate it twice: on held-out windows drawn from the *same* bearings (30% of the training recordings), and on one bearing per class that the model has never seen. Both test sets are class-balanced and the metric is macro-F1. They are the two endpoints of the ladder in Table 6, from that same run.

The optimism is +0.386, well above the CWRU average and above CWRU's worst case. It follows from drawing support and test windows from

one recording, and any benchmark organized as a few long recordings should show the same. Under a bearing-disjoint protocol the features do retain *some* transferable signal (0.505 against 0.333 chance), but of the 0.558 by which the same-recording score exceeds chance, 0.386 is optimism.

The spread is larger here than on CWRU, and one of the ten runs deserves a comment. Seed 6 failed to converge: its macro-F1 on the *training* set is 0.60 where the other nine reach 0.98–1.00, and all three of its test scores are correspondingly low. Re-running the identical configuration converged (1.00 on the training set, $L1 = 0.927$), so this is a transient optimization failure rather than a property of the split. We keep it, because the contrast is measured *within* a run and its optimism, +0.12, is the smallest of the ten (the other nine span +0.13 to +0.66). Dropping it would raise the same-recording mean from 0.891 to 0.931.

7.3. The optimism decomposes into a session term and a specimen term

The two endpoints above leave the decisive question open: between them, *two* factors change at once. The bearing-disjoint test set is a different recording session *and* a different physical bearing. CWRU cannot separate them, because each class is a single specimen. Paderborn can, because each of the numbered bearings was measured 20 times under the same condition: a model can be trained on files 1–4 of a set of bearings and tested on files 5–8 of exactly those same bearings — a different recording session of the same specimen. The three test sets are drawn as a ladder in Fig. 3(a), where the drop between consecutive points is exactly one of the two terms computed below.

Both terms are positive, and the specimen term is the larger. Simply re-measuring the *same* bearings costs +0.161 macro-F1: that much of the optimism is pure session leakage, with the specimen held fixed. Changing the specimen on top of that costs a further +0.225. The two explanations that the CWRU protocol cannot distinguish — "the model keys on the recording" and "the model keys on the individual bearing" — are therefore both real, with the specimen carrying the larger share.

The specimen term is far more variable than the session term: dropping it has a standard deviation of 0.209 against 0.046 for the session drop, and one of ten seeds (seed 3, $L3 = 0.77$) even reverses the sign. Which bearing is held out matters more than which session it was measured in, as expected if the specimen term is carried by that bearing's individual spectral signature. The ten splits between them draw on 15 of the 21 bearings (five of the six healthy,

Table 6: Decomposing the optimism, and the same decomposition with the rig tone removed. One training set (files 1–4 of five bearings per class, first 70% of windows) is evaluated on three test sets that differ by exactly one factor each: L2 moves only the recording session, L3 moves the specimen as well. Same condition, epochs and seeds as Table 5; the *no filter* column is that table's two endpoints. In the last two columns every window is masked in the frequency domain and renormalized before training, and the whole procedure is repeated from scratch: the notch removes 100 Hz and its first twenty harmonics with ±2 bins each, and the high-pass is the 2 kHz control of Table 4.

	No filter	Notch	High-pass
Spectral energy removed	—	93.3%	99.7%
L1 same file, held-out windows	0.891	0.899	0.908
L2 files 5–8, same bearings	0.730	0.772	0.776
L3 held-out bearings	0.505	0.577	0.550
Session term (L1−L2)	+0.161	+0.127	+0.132
Specimen term (L2−L3)	+0.225	+0.195	+0.226
Total optimism (L1−L3)	+0.386	+0.321	+0.358

four of the six inner-ring and six of the nine outer-ring), so what varies across seeds is a property of specimens in general rather than one unlucky draw.

Cross-specimen transfer is not zero either: $L3 = 0.505$ sits well above chance, so part of the fault physics does survive the change of specimen. Neither term rests on a handful of runs: over the ten seeds the session term is positive in all ten and the specimen term in nine, a two-sided sign test giving $p = 0.002$ and $p = 0.021$ without any normality assumption, and the 95% confidence intervals for the two means are +0.128 to +0.193 and +0.076 to +0.374. The decomposition does not depend on this network: a ResNet with a quarter of the parameters (30 099 against 118 368) and the same recipe gives $L1 = 0.883$, $L2 = 0.727$, $L3 = 0.537$ — session +0.156, specimen +0.189, the same ordering and the same total to within 0.04 (its ten runs reach a training macro-F1 of 0.98 on average, two of them 0.88).

7.4. The decomposition is not carried by the rig tone

The 100 Hz line of Section 5 belongs to the test rig rather than to any bearing, so it cannot encode a specimen's identity. Because it dominates the spectrum, however, it could carry much of the class structure on its own,

and the decomposition might then measure little more than the recovery of that one line. To test this we remove the line before training and repeat the whole decomposition with the same splits, seeds and schedule. Two filters are used: a notch at 100 Hz and its first twenty harmonics, ±2 bins each, which discards 93.3% of the spectral energy of the untapered window, and the 2 kHz high-pass of Table 4, which discards 99.7%.

Both filters leave the conclusion of the section standing (Table 6). Same-recording evaluation moves little — L1 is 0.899 and 0.908 against 0.891 unfiltered — although more than nine tenths of the energy is gone. Total optimism is +0.321 and +0.358 against +0.386. The session term is the one that moves consistently: it falls to +0.127 and +0.132 from +0.161, a paired reduction of 0.034 ± 0.012 and 0.028 ± 0.011 standard error, and it is lower in nine of the ten seeds under each filter — the exception is the run of the previous subsection that did not converge. The specimen term moves by -0.030 ± 0.018 under the notch and by $+0.001 \pm 0.023$ under the high-pass, i.e. in no direction that the two filters agree on, so we read it as unchanged. Roughly a fifth of the session leakage is therefore carried by the dominant line; the rest of it, and the whole of the specimen term, lie in the remainder of the spectrum — the bearing resonances and their sidebands, which is where the fault signature itself is.

Removing 99% of the signal energy therefore leaves the same-recording optimism at 0.32–0.36. Fig. 3 draws these numbers: in (a) the three ladders are nearly parallel, and in (b) the session bars fall under either filter while the specimen bars do not.

7.5. Where the kinematics match the labels

The Paderborn classes are defined by damage location, so here the impulse train at BPFO/BPFI is the right physics for the label. We therefore repeat the no-source few-shot experiment of Section 9 with these classes, evaluating on held-out bearings.

Three things hold in every row. Training on the real labeled samples is the highest entry, beating the pure twin in ten of ten, ten of ten and nine of ten seeds (sign test $p = 0.002$, 0.002, 0.021). The pure twin is the lowest, at 0.166 regardless of budget, even though its physics matches the class definition exactly. And the real-sample baseline does not improve systematically with k — 0.312, 0.348 and 0.321 at the three budgets — so the ceiling is set by the bearing-disjoint split rather than by sample count;

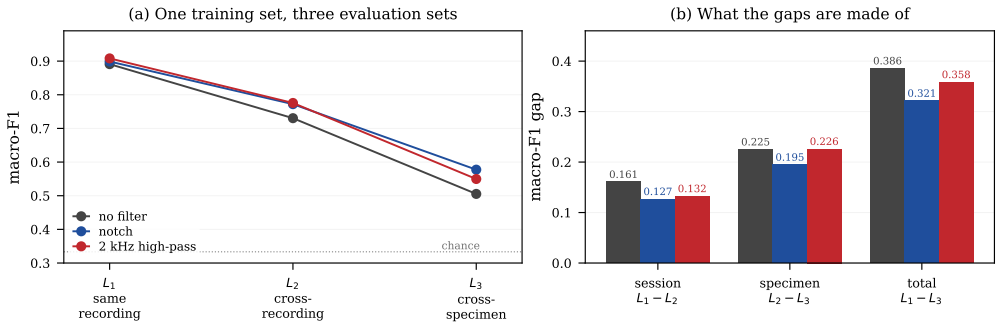


Figure 3: The decomposition of Table 6, and the same decomposition with the rig tone removed. (a) One training set, three evaluation sets that differ by one factor each; the drop between consecutive markers is exactly one of the terms plotted in (b). The three curves are the unfiltered data and the two filtered variants: removing 93.3% (notch) or 99.7% (high-pass) of the spectral energy moves any level by at most 0.07 and does not remove the decline. (b) The two terms and their sum, bars are means over ten seeds. Filtering reduces the session term from 0.161 to 0.127 and 0.132, roughly a fifth, and leaves the specimen term without a consistent direction (0.225 to 0.195 and 0.226), so the total stays in 0.32–0.36.

its standard deviation grows from 0.042 to 0.124 because the identity of the held-out bearing dominates the score.

The two envelope-based variants sit between those ends. Adding the impulse train to the target envelope moves the mean by +0.036, +0.012 and −0.004 at the three budgets (positive in 7, 6 and 4 of ten paired seeds; $p \geq 0.34$), so this benchmark does not resolve whether the physical component helps. The table therefore supports a single ordering: in every row the pure twin is below both envelope-based variants, and nothing approaches training on the real labeled samples. Not even that baseline reaches chance (0.333).

Where the physics stops hurting. On CWRU the augmentations that reproduce the target recording's spectrum are the only ones that beat the support baseline at every budget (Table 9), and adding the physical impulse train to that spectrum changes nothing at $k \leq 3$; in the synthesis-only regime, where the envelope must carry the whole class, the same physics costs 0.26 macro-F1 (Table 1). On Paderborn the same physical component does no detectable harm and may do a little good, though the difference is not resolved at this sample size. The mechanism explains the asymmetry. On CWRU the class information *is* the fingerprint, so estimating the fingerprint captures the label and physics is redundant noise. On Paderborn we measured that

Table 7: No-source few-shot augmentation on Paderborn, classes defined by damage location (healthy / inner ring / outer ring), bearing-disjoint test, macro-F1. Chance is 0.333. Ten seeds at each k; $\pm$ is the standard deviation over seeds. The last column is the same setting as Table 9 except that the class labels are, on this benchmark, consistent with the fault kinematics.

k (windows/class)	Real support	Pure twin	Fingerprint	+ physics
5	0.312 ± 0.042	0.166 ± 0.001	0.184 ± 0.029	0.220 ± 0.080
20	0.348 ± 0.084	0.166 ± 0.001	0.177 ± 0.020	0.189 ± 0.028
50	0.321 ± 0.124	0.166 ± 0.001	0.189 ± 0.034	0.185 ± 0.045

the class-mean fingerprints are identical across classes (margin +0.003, Table 4); calibrating to them therefore transfers no label information, and the only thing left that could carry the label is the impulse repetition rate — the physics.

8. A Constructive Baseline: Target-Domain Envelope Calibration

Let $\{(\mathbf{x}_i, y_i)\}_{i=1}^{kC}$ be the labeled target support set for C classes. For class c we estimate the magnitude envelope

$$A_c(f) \;=\; \frac{1}{|\mathcal{S}_c|} \sum_{i \in \mathcal{S}_c} \big| \mathcal{F}\{\mathbf{x}_i \odot \mathbf{w}\}(f) \big|, \tag{1}$$

smoothed with a moving average of width 5 bins and normalized to unit ℓ_2 norm. Every envelope number quoted in this paper (hereafter the *smoothed envelope*) uses this definition; the unsmoothed class-mean spectra of Table 4 are a different statistic and are not comparable to it. Synthetic samples are then obtained by coloring an excitation e in the frequency domain,

$$\mathbf{x}_{\text{syn}} = \mathcal{F}^{-1}\big\{\, \mathcal{F}\{e\} \cdot A_c \,\big\}, \tag{2}$$

followed by per-window z-scoring. We consider two excitations: white noise (the *fingerprint* variant) and an impulse train placed at the characteristic frequencies recomputed at the *target* shaft speed (the *fingerprint+physics* variant).

A third synthetic source, the *phase surrogate*, uses the same samples without the class mean: each of its 400 samples per class draws one of the k support windows, keeps that window's own magnitude spectrum, replaces the

phases with independent uniform draws, and inverse-transforms the result. It reproduces the recording exactly and the class only through the windows it draws from.

All three use only the k labeled support samples; no test window is involved. The envelopes are accurate at small k: estimated from $k = 5$ samples, the per-class envelope matches one estimated from the entire target load at cosine 0.995 on average (0.989 worst class). The corresponding *source-domain* envelopes match the target load only at 0.924, which is why the target-domain qualifier matters.

9. Results

9.1. Standard protocol

Table 8 reports the ten-class S0→S3 comparison, five seeds, all runs from a single uniform batch. Three observations.

The twin is worse than its own ablation. Adding the calibrated physics twin (DT__DANN) reaches macro-F1 0.914 ± 0.078, below plain adversarial adaptation (DANN, 0.986 ± 0.027) and below the physics-free augmentation ablation RAND__DANN (0.922 ± 0.098). It is also the least stable configuration.

The protocol saturates. Every configuration that uses target labels reaches 0.91–1.00, leaving almost no headroom; the differences among DANN, FP__PHYS__DANN and FP__DANN are not a meaningful demonstration of method quality in this regime. The informative comparisons are in the constrained regimes of Section 9.2.

Performance is bounded by the labeled support set. The zero-label bound SRC__PURE reaches 0.731; merely adding the $k = 5$ labeled target samples (SOURCE__ONLY) raises this to 0.823.

9.2. Constrained regimes

In the synthesis-only regime (Table 1) the physics twin is at chance (0.123) while fingerprint-colored white noise reaches 0.910; adding the physical impulse train drops this to 0.651, and removing the fault-harmonic fine structure from the envelope drops it further to 0.524.

In the no-source regime, where only the k labeled target samples are available (Table 9), the twin augmentation is worse than no augmentation at every budget (0.528, 0.811 and 0.868 against 0.841, 0.946 and 0.958). Spectral augmentation helps, most at $k = 1$: coloring white noise with the target envelope gains 0.077, 0.037 and 0.036 macro-F1 at the three budgets (17, 19 and 18 of

Table 8: Ten-class CWRU S0→S3, $k = 5$ labeled target samples per class, five seeds, single uniform batch. SRC_PURE uses no target data at all; SOURCE_ONLY and all remaining rows additionally use the labeled support set.

Method	Accuracy	macro-F1
SRC_PURE (no target data)	0.8188 ± 0.0421	0.7307 ± 0.0225
SOURCE_ONLY	0.8842 ± 0.0377	0.8228 ± 0.0686
DANN (A2 ablation)	0.9904 ± 0.0179	0.9858 ± 0.0268
RAND_DANN (A3 ablation)	0.9512 ± 0.0553	0.9223 ± 0.0977
DT_DANN (physics twin)	0.9028 ± 0.1212	0.9141 ± 0.0780
FP_PHYS_DANN (fp. + imp.)	0.9917 ± 0.0091	0.9891 ± 0.0123
fp_dann (fp. + noise)	$\mathbf{0.9998 \pm 0.0003}$	$\mathbf{0.9998 \pm 0.0004}$

20 paired runs; $p \leq 0.003$). A per-sample phase surrogate, which keeps each support window's own spectrum instead of the class mean, is statistically indistinguishable from it at every budget (+0.036, +0.010, +0.003; $p \geq 0.06$) and has the higher mean at $k = 1$. Source-domain envelopes, which need no target labels, stay below the target ones at every budget (0.844/0.923/0.979 against 0.918/0.983/0.994; 15, 17 and 15 of 20 paired runs).

Adding the impulse train to the target envelope changes the mean by at most 0.01 at $k \leq 3$ and, at $k = 5$, only through one seed whose two runs collapse to 0.014 and 0.112; apart from that seed the $k = 5$ difference is -0.004. Noise-only augmentation, which shifts and adds noise to the support windows and supplies no spectrum, is the best non-oracle entry at $k = 1$ by mean (0.964; it wins only 11 of 20 paired comparisons with the envelope) but falls to 0.239 and 0.376 in two of its twenty $k = 3$ runs, so its gain is the generic effect of enlarging a ten-window training set and is not reliable. The envelope and the surrogate, by contrast, never fall below 0.65 in any run.

At $k = 3$ and $k = 5$ the ordering follows how faithfully an augmentation reproduces the target recording: the phase surrogate, which copies each support window's spectrum exactly, is first; the smoothed envelope, which keeps the class average, next; the physics twin, which reproduces neither, last. The surrogate beats noise-only in 16 of 20 paired runs at each of $k = 3$ and $k = 5$ ($p = 0.012$). At $k = 1$ that ordering is broken by the generic effect of enlarging a ten-window training set, and noise-only augmentation is competitive. The same mechanism tells us what happens when the recording changes: under cross-recording evaluation the surrogate is the *weakest* entry

Table 9: No-source few-shot setting, target load 3, macro-F1. Each cell is mean ± SD over 20 runs (ten seeds, the whole configuration repeated twice). Augmentations add 400 samples per class generated from the k labeled support samples only; the last row draws from the target pool that is neither support nor test.

Augmentation	$k = 1$	$k = 3$	$k = 5$
none (support only)	0.841±0.059	0.946±0.045	0.958±0.031
physics twin	0.528±0.118	0.811±0.144	0.868±0.096
fingerprint + white noise	0.918±0.102	0.983±0.026	0.994±0.006
fingerprint + physics	0.927±0.042	0.982±0.016	0.897±0.279[†]
source fingerprint	0.844±0.213	0.923±0.125	0.979±0.025
phase surrogate	0.955±0.029	0.994±0.008	0.997±0.005
noise-only augmentation	0.964±0.022	0.921±0.206[‡]	0.992±0.005
true target (bound)	1.000	1.000	1.000

[†] bimodal: one seed collapses in both repetitions, to 0.014 (below the 0.10 chance level) and 0.112. [‡] likewise bimodal: two of twenty runs fall to 0.239 and 0.376, so the mean understates the usual behavior.

tested (0.777 against 0.840 for support only, Section 6).

10. Discussion and Limitations

Two benchmarks, one failure. The two datasets fail for different reasons that share a root. On CWRU the class label is confounded with the recording, so the signal a classifier learns is not the fault. On Paderborn the label *is* physically meaningful — damage location, the very quantity whose characteristic frequencies a model can compute — yet the magnitude spectrum carries none of it, because a single rig tone and its spectral leakage hold over 99% of the energy in every bearing whether healthy or damaged. A benchmark suitable for testing physics-based augmentation must satisfy both conditions at once: the label must have a physical definition, and the recording must actually expose it. Neither of the two most-used public datasets does, and establishing this belongs in benchmark validation. The CWRU benchmark has already been criticized on related grounds [26]; we add a mechanism and a quantitative control. The situation is not unique to this field: a model can generalize across hospitals while relying on a site-specific confounder invisible in the acquisition protocol [27].

The failure we diagnose lies in the *pairing* of a generator with a bench-

mark. The physical component that costs 0.26 macro-F1 when it has to carry the class alone does no detectable harm once its impulse train is layered onto the target spectral envelope: in the no-source regime it moves the mean by $+0.009$ and -0.001 at $k = 1$ and $k = 3$ (Table 9), and on Paderborn by $+0.036$, $+0.012$ and -0.004 (Table 7). The mechanism accounts for this: where the envelope already carries the fingerprint the physics is redundant, and on Paderborn, where the fingerprint carries no label, it is the only candidate carrier, though the benchmark is too noisy to resolve that into a gain. The pure twin is the worst augmentation on *both* benchmarks. Under cross-recording evaluation on CWRU the best augmentation scores 0.854 against 0.840 for training on the real labeled support samples, a margin carried by one support load.

Implications for reported numbers. Two papers reporting the same method on CWRU may differ by 0.1–0.3 macro-F1 because of how the few-shot split was drawn. Future work should state whether support and test come from the same recording, and ideally report both. The same control is cheap wherever a dataset contains several independent specimens per class: hold out specimens, not windows. Our own experience is a case in point: three seeds put the physical augmentation 0.09 above the envelope-only baseline on Paderborn, and ten seeds reduce the gap to 0.01 and remove its significance. The secondary tables in this paper are reported over ten or twenty runs rather than the conventional three, because repeated execution of *the same command* with the same seeds moved a ten-seed mean by up to 0.033 and made a three-seed standard deviation wrong by more than an order of magnitude. This is the same problem that has been documented for reinforcement learning [28] and for evaluation budgets generally [29], and it argues for reporting score *distributions* and the compute spent on them.

A benchmark validity checklist. The failures above share a structure: a benchmark can make a positive or a negative result about simulation-derived samples uninterpretable before any method is run. Table 10 states the properties a benchmark must have, the check that establishes each, the signature of its failure, and where this paper observes that signature. The first two rows cost nothing beyond a second evaluation, the third is one FFT per class, and we suggest reporting all three in the paper rather than in a supplement. The fourth row is a different matter: Section 7.4 shows that filtering the dominant line out of the data does not remove the leakage, so observability and leakage have to be checked separately rather than one standing in for the other.

Table 10: A benchmark validity checklist for simulation-based augmentation. Each row is a property a benchmark must have before a positive or a negative result about simulation-derived samples can be interpreted; the last column records where this paper observes the corresponding failure.

Dimension	Required check	Signature of failure	Observed here
Specimen independence	Each class contains independent bearings, and no specimen contributes to both support and test	One bearing reused across the split	CWRU: one bearing per class
Recording independence	Support and test from different recordings; report both numbers and the support-only baseline	A random window split of one recording, so that using any target label beats training on the support set	CWRU: +0.14–0.21 (+0.02 with sizes pooled)
Spectral separability	Within- versus between-class mean spectra, in source and target	Between-class separation no larger than within-class	Paderborn: margin +0.003
Physical observability	Fault characteristic frequencies visible above the recording's own structure	A rig tone or recording fingerprint dominates the band	Paderborn: 100 Hz line, over 99% of the energy
Temporal scale	The window spans several characteristic periods at the sampling rate	The window holds about one fault cycle	Paderborn: 1024 samples = 16 ms, one BPFO period
Synthetic validity	Synthetic and real class-conditional spectra agree, and a discriminator cannot separate them	The generator reproduces no target class	CWRU: twin macro-F1 0.123

Limitations. (i) We study two datasets and two networks of the same family, a full and a quarter-size ResNet; other architectures may behave differently. (ii) The cross-recording control on CWRU evaluates on a whole different load, which confounds the recording change with the operating-condition change; a same-load, different-session split is not available in CWRU, and we resort to Paderborn to separate the two terms of Section 7. There the separation relies on 20 repeated measurements per bearing, so it measures the session term under one rig's repeat-measurement protocol and may not transfer to datasets whose repeated recordings differ by more than a remount. (iii) Our constructive baseline is a regularizer under cross-recording

evaluation: it helps when the support recording is atypical and costs 0.02–0.03 otherwise, and the gains of Table 9 belong to a same-recording split. (iv) In the Paderborn subset we downloaded, every artificial-damage bearing is outer-ring and every run-to-failure bearing is inner-ring, so damage location and damage origin are confounded; our spectral claim is an *absence* of class structure and is unaffected, but it does not support a separate statement about outer versus inner race. (v) The Paderborn few-shot scores are low and noisy: no augmentation entry reaches chance, and with one held-out bearing per class the score is dominated by which bearing is held out (standard deviation up to 0.12). The table supports one ordering, stable at all three budgets — the pure twin below both envelope variants, the real samples above them — and nothing finer. (vi) The spectral comparison for Paderborn uses one operating condition and 21 bearings.

11. Conclusion

On CWRU cross-load diagnosis, the dominant class information is a per-bearing recording fingerprint that physics cannot predict. Simulation-based sample supplementation is therefore structurally ineffective on that benchmark, and the standard few-shot protocol hides this by leaking the fingerprint from support to test, inflating reported macro-F1 by 0.14–0.21 (0.29 in the hardest direction). In the no-source regime the augmentations that work reproduce the target recording's spectrum — a class-mean envelope or a per-sample phase surrogate — and beat training on the labeled support at every budget; the physics twin is worse than doing nothing and adds nothing to the envelope.

Paderborn tests whether this picture is specific to CWRU. On Paderborn, where several independent bearings share each class, class-mean magnitude spectra of healthy, inner-ring and outer-ring bearings are indistinguishable (margin +0.003 after the rig tone is filtered out, against 0.626 on CWRU), so CWRU's class separation is a property of individual specimens rather than of the classes. Same-recording evaluation there is still optimistic, by +0.386, of which +0.161 is attributable to the recording session alone and +0.225 to the change of specimen. Removing the rig tone, which holds 93% of the spectral energy, leaves +0.321, still dominated by the specimen term. There the labels follow the fault kinematics and the fingerprints carry none of them, and the physical component that costs 0.26 macro-F1 on CWRU does no detectable harm. The pure twin is the worst augmentation on both benchmarks, and

under cross-recording or bearing-disjoint evaluation no augmentation beats training on the real labeled samples by more than 0.015. Table 10 lists the six checks we would run before trusting a result, positive or negative, on a new benchmark.

CRediT authorship contribution statement

Fangyu Zhang: Conceptualization, Methodology, Software, Formal analysis, Investigation, Resources, Writing – original draft, Writing – review & editing, Project administration. **Julong He:** Supervision, Validation, Data curation, Visualization.

Declaration of competing interest

The authors declare that they have no known competing financial interests or personal relationships that could have appeared to influence the work reported in this paper.

Funding

This research did not receive any specific grant from funding agencies in the public, commercial, or not-for-profit sectors.

Data availability

The CWRU 12 kHz drive-end and Paderborn KAt datasets analysed here are public. The code and the per-table artifacts that reproduce every number in this paper are available at `https://github.com/qy321666/WSlB` (release `v2.4`) and are also provided as supplementary material to this submission.

Declaration of generative AI and AI-assisted technologies in the writing process

During the preparation of this work the authors used ClawsGO Science Agent in order to improve the language of the manuscript. After using this tool, the authors reviewed and edited the content as needed and take full responsibility for the content of the publication.

References

[1] Y. Ganin, E. Ustinova, H. Ajakan, P. Germain, H. Larochelle, F. Laviolette, M. Marchand, V. Lempitsky, Domain-adversarial training of neural networks, Journal of Machine Learning Research 17 (59) (2016) 1–35.

[2] B. Sun, J. Feng, K. Saenko, Return of frustratingly easy domain adaptation, Proceedings of the AAAI Conference on Artificial Intelligence 30 (1) (2016). doi:10.1609/aaai.v30i1.10306.

[3] Z. Ming, B. Tang, L. Deng, X. Zhang, Z. Li, Physics-enhanced simulation-to-measurement translation for rolling bearing fault diagnosis under limited samples, Engineering Applications of Artificial Intelligence 164 (2026) 113279. doi:10.1016/j.engappai.2025.113279.

[4] X. Xiao, Y. He, J. Su, K. Hu, Domain adaptation from simulation to reality: A GAN- and MK-MMD-based transfer learning approach for bearing fault diagnosis, Applied Sciences 16 (3) (2026) 1407. doi:10.3390/app16031407.

[5] X. Liu, Y. Jin, F. Yang, Y. Kang, L. Bo, Twin data multimode collaborative transfer learning for bearing failure diagnosis, Engineering Applications of Artificial Intelligence 137 (2024) 109142. doi:10.1016/j.engappai.2024.109142.

[6] S. Shao, P. Wang, R. Yan, Generative adversarial networks for data augmentation in machine fault diagnosis, Computers in Industry 106 (2019) 85–93. doi:10.1016/j.compind.2019.01.001.

[7] G. E. Karniadakis, I. G. Kevrekidis, L. Lu, P. Perdikaris, S. Wang, L. Yang, Physics-informed machine learning, Nature Reviews Physics 3 (6) (2021) 422–440. doi:10.1038/s42254-021-00314-5.

[8] F. Tao, H. Zhang, A. Liu, A. Y. C. Nee, Digital twin in industry: State-of-the-art, IEEE Transactions on Industrial Informatics 15 (4) (2019) 2405–2415. doi:10.1109/TII.2018.2873186.

[9] Y. Xu, Q. Jiang, S. Li, Z. Zhao, B. Sun, G. Q. Huang, Digital twin-driven discriminative graph learning networks for cross-domain bearing fault recognition, Computers & Industrial Engineering 193 (2024) 110292. doi:10.1016/j.cie.2024.110292.

[10] Y. Zhang, J. C. Ji, Z. Ren, Q. Ni, F. Gu, K. Feng, Digital twin-driven partial domain adaptation network for intelligent fault diagnosis of rolling bearing, Reliability Engineering & System Safety 234 (2023) 109186. doi:10.1016/j.ress.2023.109186.

[11] A. Kumar, R. Kumar, J. Xiang, Z. Qiao, Y. Zhou, H. Shao, Digital twin-assisted ai framework based on domain adaptation for bearing defect diagnosis in the centrifugal pump, Measurement 235 (2024) 115013. doi:10.1016/j.measurement.2024.115013.

[12] Y. Lei, B. Yang, X. Jiang, F. Jia, N. Li, A. K. Nandi, Applications of machine learning to machine fault diagnosis: A review and roadmap, Mechanical Systems and Signal Processing 138 (2020) 106587. doi:10.1016/j.ymssp.2019.106587.

[13] Z. Zhao, T. Li, J. Wu, C. Sun, S. Wang, R. Yan, X. Chen, Deep learning algorithms for rotating machinery intelligent diagnosis: An open source benchmark study, ISA Transactions 107 (2020) 224–255. doi:10.1016/j.isatra.2020.08.010.

[14] M. Wang, W. Deng, Deep visual domain adaptation: A survey, Neurocomputing 312 (2018) 135–153. doi:10.1016/j.neucom.2018.05.083.

[15] Z. Zhao, Q. Zhang, X. Yu, C. Sun, S. Wang, R. Yan, X. Chen, Applications of unsupervised deep transfer learning to intelligent fault diagnosis: A survey and comparative study, IEEE Transactions on Instrumentation and Measurement 70 (2021) 1–28. doi:10.1109/TIM.2021.3116309.

[16] S. Wang, D. Wang, D. Kong, J. Wang, W. Li, Few-shot rolling bearing fault diagnosis with metric-based meta learning, Sensors 20 (22) (2020) 6437. doi:10.3390/s20226437.

[17] J. Jiao, M. Zhao, J. Lin, K. Liang, A comprehensive review on convolutional neural network in machine fault diagnosis, Neurocomputing 417 (2020) 36–63. doi:10.1016/j.neucom.2020.07.088.

[18] J. Tobin, R. Fong, A. Ray, J. Schneider, W. Zaremba, P. Abbeel, Domain randomization for transferring deep neural networks from simulation to the real world, in: IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS), 2017, pp. 23–30. doi:10.1109/IROS.2017.8202133.

[19] J. Theiler, S. Eubank, A. Longtin, B. Galdrikian, J. Doyne Farmer, Testing for nonlinearity in time series: the method of surrogate data, Physica D: Nonlinear Phenomena 58 (1-4) (1992) 77–94. doi:10.1016/0167-2789(92)90102-S.

[20] T. Schreiber, A. Schmitz, Surrogate time series, Physica D: Nonlinear Phenomena 142 (3-4) (2000) 346–382. doi:10.1016/S0167-2789(00)00043-9.

[21] W. A. Smith, R. B. Randall, Rolling element bearing diagnostics using the case western reserve university data: A benchmark study, Mechanical Systems and Signal Processing 64 (2015) 100–131.

[22] R. B. Randall, J. Antoni, Rolling element bearing diagnostics—a tutorial, Mechanical Systems and Signal Processing 25 (2) (2011) 485–520.

[23] R. Geirhos, J.-H. Jacobsen, C. Michaelis, R. Zemel, W. Brendel, M. Bethge, F. A. Wichmann, Shortcut learning in deep neural networks, Nature Machine Intelligence 2 (11) (2020) 665–673. doi:10.1038/s42256-020-00257-z.

[24] S. Kapoor, A. Narayanan, Leakage and the reproducibility crisis in machine-learning-based science, Patterns 4 (9) (2023) 100804. doi:10.1016/j.patter.2023.100804.

[25] C. Lessmeier, J. K. Kimotho, D. Zimmer, W. Sextro, Condition monitoring of bearing damage in electromechanical drive systems by using motor current signals of electric motors: A benchmark data set for data-driven classification, PHM Society European Conference 3 (1) (2016). doi:10.36001/phme.2016.v3i1.1577.

[26] J. Hendriks, P. Dumond, D. A. Knox, Towards better benchmarking using the cwru bearing fault dataset, Mechanical Systems and Signal Processing 169 (2022) 108732. doi:10.1016/j.ymssp.2021.108732.

[27] J. R. Zech, M. A. Badgeley, M. Liu, A. B. Costa, J. J. Titano, E. K. Oermann, Variable generalization performance of a deep learning model to detect pneumonia in chest radiographs: A cross-sectional study, PLOS Medicine 15 (11) (2018) e1002683. doi:10.1371/journal.pmed.1002683.

[28] P. Henderson, R. Islam, P. Bachman, J. Pineau, D. Precup, D. Meger, Deep reinforcement learning that matters, Proceedings of the AAAI Conference on Artificial Intelligence 32 (1) (2018). doi:10.1609/aaai.v32i1.11694.

[29] J. Dodge, S. Gururangan, D. Card, R. Schwartz, N. A. Smith, Show your work: Improved reporting of experimental results, in: Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP), 2019, pp. 2185–2194. doi:10.18653/v1/D19-1224.